

HIGH PERFORMANCE PARALLEL PROGRAMMING (CS61064)

Soumyajit Dey
CSE, IIT Kharagpur

Machine Learning

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

- Learning Rules/Functions from Examples
- Deep Learning (a branch of ML) has gained significant success over the past few years.
- Advancements in Computer Architecture and GPU Programming have allowed ideas developed in the 80s to be fully realized.

Examples

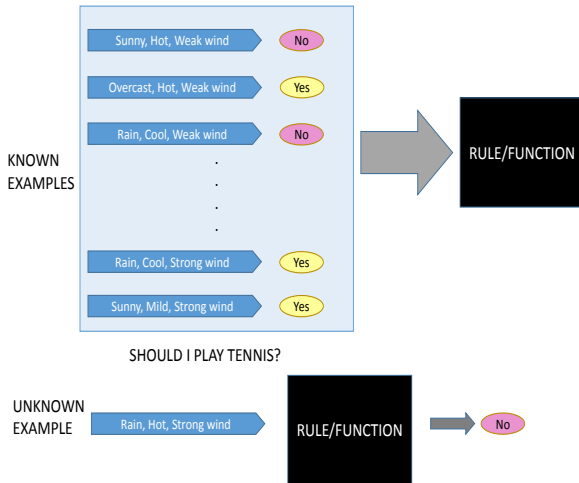


Figure: Decision Making

Examples

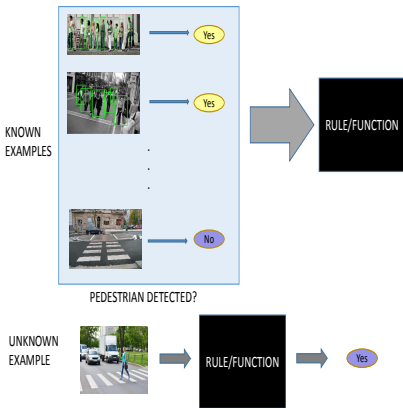


Figure: Pedestrian Detection

ML Problem Characteristics

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

- Characterize each example by a vector of real or boolean values (feature vector)
- Associate a label/target with each example (Supervised Learning).
- Labels may be discrete values (Classification) or continuous values (Regression)
- A set of feature vectors and labels constitute a training data set.
- Using the training dataset and some supervised learning algorithm a predictive model is trained.

Supervised Learning Workflow

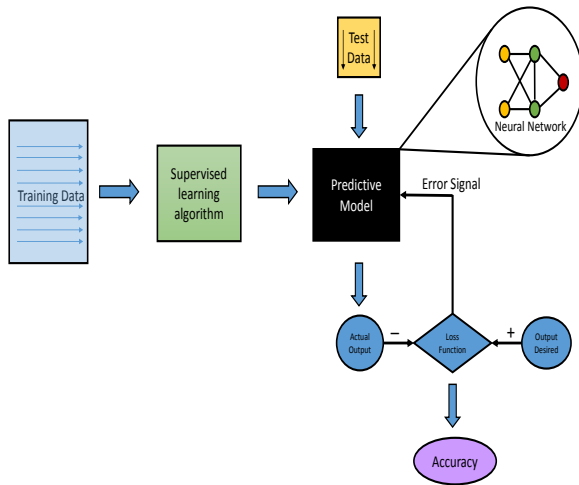


Figure: Pedestrian Detection

Supervised Learning Algorithm

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

- Characterize the model to be learned by some parameter θ
- Define a loss function between predicted output and actual output. (function of θ and inputs)
- Update θ so that the loss function is minimized.
- The more the loss function is closer to the minima, the more closer is the predicted output to the actual output

Neural Networks

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

- Neural networks have gained popularity as parametric learning methodologies, due to advancement of Deep Learning.
- The methodology is derived directly from the working model of the human brain.
- Can be depicted as a network of weights associating input values to output values.

Supervised Learning Workflow




Figure: Pedestrian Detection

Building Block of a NN




Figure: Perceptron


- The yellow nodes are inputs while the grey node is a neuron.
- The edges(synapses) have weights)
- The incoming value to the neuron is $f = \sum w_i x_i$
- The outgoing value is a nonlinear function of f i.e.
 $o = \sigma(f)$

Activation Functions

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)


Soumyajit Dey
CSE, IIT
Kharagpur

SIGMOID




$$y = \frac{1}{1 + e^{-x}}$$

HYPERBOLIC
TANGENT



$$y = \frac{1 - e^{-2x}}{1 + e^{-2x}}$$

RECTIFIED
LINEAR UNIT



$$y = \max(0, x)$$

Figure: Linear/Non-linear functions

Multilayer Perceptron/ FeedForward Network

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur




Figure: Classification

Multilayer Perceptron/ FeedForward Network

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur




Figure: Regression

Note

- The objective of this course is to get you acquainted with the computation involved while training and testing a neural network.
- We shall not discuss core ML principles which should be followed while designing neural networks for various problem domains.

Neural Networks

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

- Each neuron of a layer accumulates a weighted sum of the inputs from the previous layer.
- Each neuron applies an activation function to its input and propagates the output to a neuron of the next layer..
- This results in a series of linear and non-linear transformations from the input layer to the output layer.
- The predicted output value for every input example is a function of the input feature values and the weights and activation in the network

Feedforward NN Forward Propagation

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

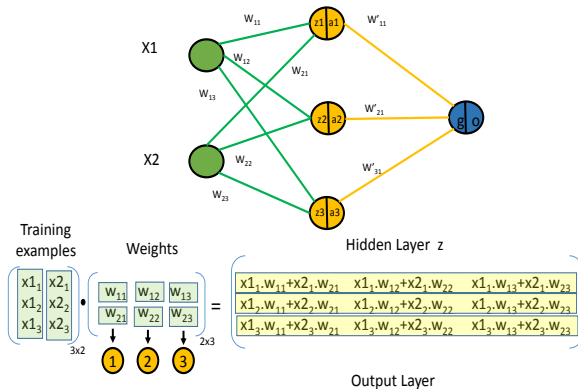


Figure: Feedforward propagation

Feedforward NN Forward Propagation

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

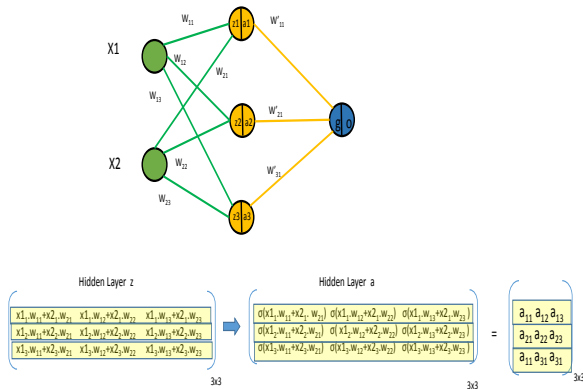


Figure: Feedforward propagation

Feedforward NN Forward Propagation

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

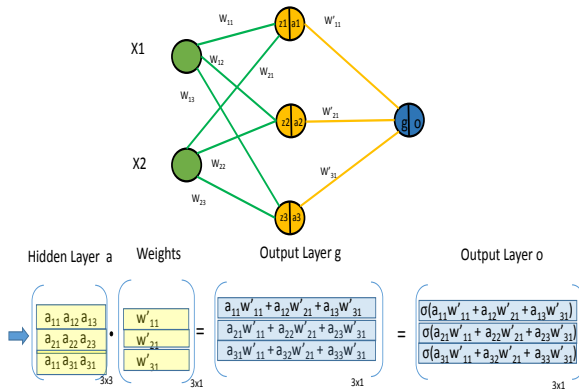



Figure: Feedforward propagation

Neural Networks



$$Z = XW \longrightarrow A = f(Z) \longrightarrow G = AW' \longrightarrow O = f(G)$$

Figure: Feedforward propagation

The feedforward propagation can therefore be expressed as a series of matrix computations and element-wise non-linear transformations which involves scope for GPU parallelization.

Neural Networks

- Given the structure and weights of a neural network, we now know how to compute the predicted output value for an input example.
- The structure of the network i.e. the number of layers and number of neurons per layer are referred as hyperparameters (to be decided by the user).
- The weights are the actual parameters (θ) which will be learned during the course of training.
- Training involves the minimization of a cost/loss function.

Loss Function

- Define a loss function $J(w) = \frac{1}{2} \sum_{i=1}^n (y_i - o_i)^2$ where y_i and o_i are the actual outputs and predicted outputs of input example i respectively.
- Recall the feedforward propagation equations.
 $\mathbf{Z} = \mathbf{XW}$, $\mathbf{A} = f(\mathbf{Z})$, $\mathbf{G} = \mathbf{AW}'$, $\mathbf{O} = f(\mathbf{G})$
- Therefore, $J = \sum \frac{1}{2} (\mathbf{Y} - f(f(\mathbf{XW})\mathbf{W}'))^2$ where the summation operation is over the elements of the column vector obtained from the loss function.
-

Minimizing Loss Function

- Find values of \mathbf{W} and \mathbf{W}' so that $J(\mathbf{w})$ is minimized.
- Compute $\frac{\partial J}{\partial \mathbf{W}}$ and $\frac{\partial J}{\partial \mathbf{W}'}$
- Instead of setting the partial derivatives to zero and finding a solution, we perform numerical gradient descent.
- Update the weights \mathbf{W} and \mathbf{W}' in the direction of the steepest gradient descent

Gradient Descent




Figure: Obtaining minimum J

Gradient Descent

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur




Figure: Obtaining minimum J

Training Using Backpropagation

- Perform feedforward propagation to obtain predicted output values for each input example.
- Compute loss function $J(w)$
- Compute $\frac{\partial J}{\partial W}$ for each weight matrix
- Update weights of every weight matrix W with the gradient.
- Repeat steps 1-3 until there is no change in gradient.

Computing Partial Derivatives w.r.t a Matrix


$$\frac{\partial J}{\partial W} = \begin{bmatrix} \frac{\partial J}{w_{11}} & \frac{\partial J}{w_{12}} & \frac{\partial J}{w_{13}} & \cdots & \frac{\partial J}{w_{1n}} \\ \frac{\partial J}{w_{21}} & \frac{\partial J}{w_{22}} & \frac{\partial J}{w_{23}} & \cdots & \frac{\partial J}{w_{2n}} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \frac{\partial J}{w_{d1}} & \frac{\partial J}{w_{d2}} & \frac{\partial J}{w_{d3}} & \cdots & \frac{\partial J}{w_{dn}} \end{bmatrix}$$

The dimensions of the weight matrix and its gradients will be the same.

Neural Networks

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur



$$Z = XW \longrightarrow A = f(Z) \longrightarrow G = AW' \longrightarrow O = f(G)$$

Figure: Feedforward propagation

We first compute $\frac{\partial J}{\partial W'}$

Chain Rule

$$\begin{aligned}\frac{\partial J}{\partial \mathbf{W}'} &= \frac{\partial}{\partial \mathbf{W}'} \sum \frac{1}{2} (\mathbf{Y} - \mathbf{O})^2 = - \sum (\mathbf{Y} - \mathbf{O}) \frac{\partial \mathbf{O}}{\partial \mathbf{G}} \frac{\partial \mathbf{G}}{\partial \mathbf{W}'} \\ &= - \sum (\mathbf{Y} - \mathbf{O}) f'(\mathbf{G}) \frac{\partial \mathbf{G}}{\partial \mathbf{W}'}\end{aligned}$$

Consider input example 1 and one weight say w'_{11}

The derivative is

$$\begin{aligned}&-(y_1 - o_1) f'(g_1) \frac{\partial}{\partial w'_{11}} (w'_{11} a_{11} + w'_{21} a_{12} + w'_{31} a_{13}) \\ &= \delta_1^1 a_{11}\end{aligned}$$

For input example 2, the gradient would be $= \delta_2^1 a_{21}$

For input example 3, the gradient would be $= \delta_3^1 a_{31}$

Computing Partial Derivatives w.r.t a Matrix

$$\frac{\partial J}{\partial W'} = \begin{bmatrix} \frac{\partial J}{w'_{11}} \\ \frac{\partial J}{w'_{21}} \\ \frac{\partial J}{w'_{31}} \end{bmatrix} = \begin{bmatrix} \delta_1^1 a_{11} + \delta_2^1 a_{21} + \delta_3^1 a_{31} \\ \delta_1^1 a_{12} + \delta_2^1 a_{22} + \delta_3^1 a_{32} \\ \delta_1^1 a_{13} + \delta_2^1 a_{23} + \delta_3^1 a_{33} \end{bmatrix}$$

This can be expressed as $\mathbf{A}^T \delta^1$

Computing gradient w.r.t W




Figure: Feedforward propagation

In a similar fashion we compute $\frac{\partial J}{\partial W}$

Chain Rule

$$\begin{aligned}\frac{\partial J}{\partial \mathbf{W}} &= -\sum (Y - O) \frac{\partial O}{\partial \mathbf{G}} \frac{\partial \mathbf{G}}{\partial \mathbf{W}} \\&= -\sum (Y - O) f'(G) \frac{\partial \mathbf{G}}{\partial \mathbf{A}} \frac{\partial \mathbf{A}}{\partial \mathbf{W}} \\&= \sum \delta^1 \mathbf{W}'^T \frac{\partial \mathbf{A}}{\partial \mathbf{W}} \\&= \sum \delta^1 \mathbf{W}'^T \frac{\partial \mathbf{A}}{\partial \mathbf{Z}} \frac{\partial \mathbf{Z}}{\partial \mathbf{W}} \\&= \sum \delta^1 \mathbf{W}'^T f'(\mathbf{Z}) \frac{\partial \mathbf{Z}}{\partial \mathbf{W}} \\&= \mathbf{X}^T \delta^1 \mathbf{W}'^T f'(\mathbf{Z}) \text{ (Derive!)} \\&= \mathbf{X}^T \delta^2\end{aligned}$$

Backward Propagation Delta Rule




Figure: Backward propagation

Summary

- Feedforward propagation can be performed by a series of linear and non linear transformations involving matrix operations starting from the input layer.
- Backpropagation also involves a series of linear and non linear transformations involving matrix operations starting from the output layer.
- Each operation and transformation exhibits parallelism and scope for optimizations using a GPU.

Building a DL Library

A DL Library should have support for the following

- Provide constructs for specifying a network.
- Provide efficient routines for feedforward, backpropagation and gradient computation.
- Provide routines for training and testing.
- Should support parallel and distributed processing for the computation passes.

DL Libraries like Tensorflow and Theano use a Computational Graph Abstraction for encoding a neural network.

Computational Graph

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur






Figure: MLP vs CG

Computational Graph

A graph that denotes the functional description of the required computation.

- A node with no incoming edge is a tensor, matrix, vector or scalar value.
- A node with an incoming edge is a function of the edge's tail node. computation.
- An edge represents a data dependency between nodes.
- A node knows how to compute its value and the value of its derivative w.r.t each incoming edges's tail node.

Computational Graph



J is the loss function

$$Z = XW \quad A = f(Z) \quad G = AW' \quad O = f(G)$$




Figure: MLP vs CG

Computations for a CG

- **Forward Computation:** Loop over each node in topological order and compute the value of the node given its inputs.
- **Backward Computation** Loop over each node in reverse topological order and compute the derivative of the final goal node with respect to each incoming edge's tail node.

Backpropagation: Gradient w.r.t W'

HIGH
PERFORMANCE
PARALLEL
PROGRAMMING
(CS61064)

Soumyajit Dey
CSE, IIT
Kharagpur

$$\frac{\partial J}{\partial W} = \frac{\partial J}{\partial O} \frac{\partial O}{\partial G} \frac{\partial G}{\partial A} \frac{\partial A}{\partial Z} \frac{\partial Z}{\partial W}$$

J is the loss function

$Z = XW$ $A = f(Z)$ $G = AW'$ $O = f(G)$




Figure: Computing gradients on CG

Backpropagation: Gradient w.r.t W

$$\frac{\partial J}{\partial W'} = \frac{\partial J}{\partial O} \frac{\partial O}{\partial G} \frac{\partial G}{\partial W'}$$

J is the loss function

$Z = XW$ $A = f(Z)$ $G = AW'$ $O = f(G)$




Figure: Computing gradients on CG