

Lecture 05-06

Example -2

```
# include <stdio.h>
# include <omp.h>
int main ( int argc, char *argv[] ) {
    int id;
    double wtime;
    printf ( "Number of processors available = %d\n", omp_get_num_procs ( ) );
    printf ( "Number of threads =          %d\n", omp_get_max_threads ( ) );
    wtime = omp_get_wtime ( );
    printf ( "OUTSIDE the parallel region.\n" );

    id = omp_get_thread_num ( );
    printf ( "HELLO from process %d\n Going INSIDE the parallel region:\n ", id );

    # pragma omp parallel \
    private ( id ) {
        id = omp_get_thread_num ( );
        printf ( " Hello from process %d\n", id );
    }
    wtime = omp_get_wtime ( ) - wtime;

    printf ( "Back OUTSIDE the parallel region.\nNormal end of execution.\nElapsed wall clock time = %f\n", wtime );
    return 0;
}
```

RECAP

Example -2

RECAP

Initialization:

```
export OMP_NUM_THREADS=16
```

Compilation:

```
g++ -fopenmp example.c
```

Execution:

```
./a.out
```

Example -4

RECAP

```
#include <stdio.h>
#include <unistd.h>
#include <omp.h>
int main()
{
    int i,j,n,m,temp,a[100][100];
    n=m=7;
    #pragma omp parallel
    {
        for(i=0;i<=n*m-1;i++) {
            temp=i/m+1;
            j=i%m+1;
            sleep(1);
            a[temp][j]=temp+100*(j-1);
        }
    }
    for(i=0;i<=n*m-1;i++) {
        temp=i/m+1;
        j=i%m+1;
        if(i%m==0) printf("\n");
        printf("%d\t",a[temp][j]);
    }
    printf("\n");
    return 0;
}
```

\$./a.out						
1	0	201	0	401	0	601
0	102	0	302	0	502	0
3	0	203	0	403	0	603
0	104	0	304	0	504	0
5	0	205	0	405	0	605
0	106	0	0	0	506	606
7	0	207	0	407	0	607

Example -4

```
#include <stdio.h>
#include <unistd.h>
#include <omp.h>
int main()
{
    int i,j,n,m,temp,a[100][100];
    n=m=7;
    #pragma omp parallel private (temp, j)
    {
        for(i=0;i<=n*m-1;i++) {
            temp=i/m+1;
            j=i%m+1;
            sleep(1);
            a[temp][j]=temp+100*(j-1);
        }
    }
    for(i=0;i<=n*m-1;i++) {
        temp=i/m+1;
        j=i%m+1;
        if(i%m==0) printf("\n");
        printf("%d\t",a[temp][j]);
    }
    printf("\n");
    return 0;
}
```

RECAP

\$./a.out

1	101	201	301	401	501	601
2	102	202	302	402	502	602
3	103	203	303	403	503	603
4	104	204	304	404	504	604
5	105	205	305	405	505	605
6	106	206	306	406	506	606
7	107	207	307	407	507	607

Example -5

```
#include <stdio.h>
#include <omp.h>
int main()
{
    int i,x=0, a[10],b[10];
    for(i=0;i<10;i++) {
        a[i]=i; b[i]=10-i;
    }
    #pragma omp parallel
    {
        #pragma omp for reduction (+:x)
        for(i=0;i<10;i++) {
            x = x + a[i]*b[i];
        }
    }
    printf("%d\n",x);
    return 0;
}
```

RECAP

\$./a.out

165

\$

Work-sharing Constructs

- **sections Construct**

```

— #pragma omp sections [clause[[,] clause] ...] new-line
{
    [#pragma omp section new-line
     structured-block]
    [#pragma omp section new-line
     structured-block ]
    ...
}

```

- **Clause**

- **private**(*variable-list*)
- **firstprivate**(*variable-list*)
- **lastprivate**(*variable-list*)
- **reduction**(*operator: variable-list*)
- **nowait**

Work-sharing Constructs

- **single Construct**

```

— #pragma omp single [clause[[,] clause] ...] new-line
   structured-block

```

- **Clause**

- **private**(*variable-list*)
- **firstprivate**(*variable-list*)
- **copyprivate**(*variable-list*)
- **nowait**

The SINGLE construct allows code that is serial in nature to be executed inside a parallel region. The thread executing the code will be the first to reach the directive in the code. It doesn't have to be the master thread. All other threads proceed to the end of the structured block where there is an implicit synchronization.

Master construct

```
#pragma omp master
    structured-block
```

Same as single nowait but only for master thread

Allowed Combinations

Clause	PARALLEL	DO/for	SECTION	SINGLE	WORKSHARE	PARALLEL	PARALLEL	PARALLEL
		r			E	DO/for	SECTIONS	WORKSHARE
IF	OK					OK	OK	OK
PRIVATE	OK	OK	OK	OK		OK	OK	OK
SHARED	OK	OK				OK	OK	OK
DEFAULT	OK					OK	OK	OK
FIRSTPRIVATE	OK	OK	OK	OK		OK	OK	OK
LASTPRIVATE		OK	OK			OK	OK	
REDUCTION	OK	OK	OK			OK	OK	OK
COPYIN	OK					OK	OK	OK
SCHEDULE		OK				OK		
ORDERED		OK				OK		
NOWAIT		OK	OK	OK	OK			

Synchronization

- CRITICAL: Mutual Exclusion
- ATOMIC: Atomic Update
- BARRIER: Barrier Synchronization

Critical construct


```
#pragma omp critical [name]  
    structured-block
```

Example

```
#pragma omp parallel  
{  
    #pragma omp critical(long_critical_name)  
        doSomeCriticalWork_1();  
    #pragma omp critical  
        doSomeCriticalWork_2();  
    #pragma omp critical  
        doSomeCriticalWork_3();  
}
```

Critical construct

➞ Only one thread at time in the critical section



Example -6

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
int main()
{
    int x=0, size=12;

    x=x+1;
    printf("%d\n",x);
    return 0;
}
```

```
$ gcc -Wall example2.c
$ ./a.out
1
$
```

Example -6

```
#include <stdio.h>
#include <omp.h>

int main()
{
    int x=0, size=12;

    omp_set_num_threads(size);
    #pragma omp parallel shared(x)
    {
        x=x+1;
    }
    printf("%d\n",x);
    return 0;
}
```

```
$ gcc -Wall -fopenmp example2_openMP.c
$ ./a.out
10
$ ./a.out
12
$ ./a.out
12
$ ./a.out
12
$ ./a.out
12
$ ./a.out
11
$
```

Example -6

```
#include <stdio.h>
#include <omp.h>

int main()
{
    int x=0, size=12;
    omp_set_num_threads(size);
    #pragma omp parallel shared(x)
    {
        #pragma omp critical
        {
            x=x+1;
        }
    }
    printf("%d\n",x);
    return 0;
}
```

```
$ gcc -Wall -fopenmp example2_openMP2.c
$ ./a.out
12
$ ./a.out
12
$ ./a.out
12
$ ./a.out
12
$ ./a.out
12
$ ./a.out
12
$
```


Important Notes

- ➔ Loop indexes are automatically PRIVATE
- ➔ Everything “local or temporary” should be PRIVATE (or FIRSTPRIVATE or LASTPRIVATE if its value is used outside the loop, before or after respectively)
- ➔ Everything “persistent” and/or used for different values of the loop index should be SHARED
- ➔ a SHARED variable that is not an array accessed with the loop indexes, should be written only in a CRITICAL region (serialize, so it's slow)
- ➔ If you are using CRITICAL, see if REDUCTION is an option (maybe changing the math a little bit)

Barrier construct

`#pragma omp barrier //Threads wait until all threads reach this point`

Example (waiting for the master to come)

```
int counter = 0;
#pragma omp parallel
{
    #pragma omp master
    counter = 1;
    #pragma omp barrier
    printf("%d\n", counter);
}
```

Be careful not to cause deadlock:
No barrier inside of *critical*, *master*, *sections*, *single*!

Atomic construct

```
#pragma omp atomic \  
    [read | write | update | capture]  
    expression-stmt
```

```
#pragma omp atomic capture  
    structured-block
```

Work out - 1

- Read an integer number by master thread only
- Fork 8 threads
- Generate i^{th} prime number (i is the thread number)
- Multiply with the input number
- Output the result as ordered by the thread number (thread 0, 1, ...)

Combined Parallel Work-sharing Constructs

- **parallel for Construct**
 - `#pragma omp parallel for` [*clause* [, *clause*] ...] *new-line*
 - *for-loop*

Combined Parallel Work-sharing Constructs

- **parallel sections Construct**

```
#pragma omp parallel sections [clause [, clause] ...] new-line
{
    [#pragma omp section new-line]
    structured-block
    [#pragma omp section new-line]
    structured-block ]
...
}
```

Loop construct: Scheduling

```
#pragma omp for schedule(kind[, chunk_size])  
for-loops
```

Static

- iterations are divided into chunks of size `chunk_size`
- the chunks are assigned to the threads in a round-robin fashion
- must be reproducible within the same parallel region

Dynamic

- iterations are divided into chunks of size `chunk_size`
- the chunks are assigned to the threads as they request them
- the default `chunk_size` is 1

Guided

- iterations are divided into chunks of decreasing size
- the chunks are assigned to the threads as they request them
- `chunk_size` controls the minimum size of the chunks

Run-time

- controlled by environment variables

Work out - 2

- Write a serial program to output the prime numbers occurring between 1 and 131072. Report the time required to compute the i^{th} prime.
- Convert it to a OpenMP code. Report the percentage of improvement for each prime number over serial program using 4 and 16 cores.

Example 8: Matrix Multiplication

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

#define NRA 62      /* number of rows in matrix A */
#define NCA 15      /* number of columns in matrix A */
#define NCB 7       /* number of columns in matrix B */

int main (int argc, char *argv[])
{
    int  tid, nthreads, i, j, k, chunk;
    double a[NRA][NCA],
           b[NCA][NCB],
           c[NRA][NCB];    /* result matrix*/

    chunk = 10;            /* set loop iteration chunk size */
```

Example 8: Matrix Multiplication

```
#pragma omp parallel shared(a,b,c,nthreads,chunk) private(tid,i,j,k)
{
    tid = omp_get_thread_num();
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Starting matrix multiple example with %d threads\n",nthreads);
        printf("Initializing matrices...\n");
    }

    #pragma omp for schedule (static, chunk)
    for (i=0; i<NRA; i++)
        for (j=0; j<NCA; j++)
            a[i][j]= i+j;
    #pragma omp for schedule (static, chunk)
    for (i=0; i<NCA; i++)
        for (j=0; j<NCB; j++)
            b[i][j]= i*j;
    #pragma omp for schedule (static, chunk)
    for (i=0; i<NRA; i++)
        for (j=0; j<NCB; j++)
            c[i][j]= 0;
```

Example 8: Matrix Multiplication

```

/** Do matrix multiply sharing iterations on outer loop */
/** Display who does which iterations for demonstration purposes */
printf("Thread %d starting matrix multiply...\n",tid);
#pragma omp for schedule (static, chunk)
for (i=0; i<NRA; i++) {
    printf("Thread=%d did row=%d\n",tid,i);
    for(j=0; j<NCB; j++)
        for (k=0; k<NCA; k++)
            c[i][j] += a[i][k] * b[k][j];
}
} /** End of parallel region */

/** Print results */
printf("*****\n");
printf("Result Matrix:\n");
for (i=0; i<NRA; i++) {
    for (j=0; j<NCB; j++)
        printf("%6.2f  ", c[i][j]);
    printf("\n");
}
printf("*****\n");
printf ("Done.\n");
}

```

Question???

- Is the static kind optimum?

Loop construct: Nested Loops

```
#pragma omp parallel
{
    #pragma omp for
    for(int ii = 0; ii < n; ii++) {
        #pragma omp for
        for(int jj = 0; jj < m; jj++) {
            A[ii][jj] = ii*m + jj;
        }
    }
}
```

Collapse clause

- ➔ Available in OpenMP 3.0 and later
- ➔ Clause for the PARALLEL DO directive
- ➔ Cause an automatic “collapse” (merge) of the loops, and thus automatic parallelization of inner loops too
- ➔ Most things taken care automatically, but user have to be careful

Loop construct: Nested Loops

```
#pragma omp parallel
{
    #pragma omp for collapse(2)
    for(int ii = 0; ii < n; ii++) {
        for(int jj = 0; jj < m; jj++) {
            A[ii][jj] = ii*m + jj;
        }
    }
}
```

The collapsed loops must be perfectly nested.

Collapse Clause

- **Purpose**
 - Specifying the COLLAPSE clause allows you to parallelize multiple loops in a nest without introducing nested parallelism.
- **Rules**
 - Only one collapse clause is allowed on a work sharing DO or PARALLEL DO directive
 - The specified number of loops must be present lexically. That is, none of the loops can be in a called subroutine.
 - The loops must form a rectangular iteration space and the bounds and stride of each loop must be invariant over all the loops.
 - If the loop indices are of different size, the index with the largest size will be used for the collapsed loop.
 - The loops must be perfectly nested; that is, there is no intervening code nor any OpenMP directive between the loops which are collapsed.
 - The associated do-loops must be structured blocks. Their execution must not be terminated by an EXIT statement.
 - If multiple loops are associated to the loop construct, only an iteration of the innermost associated loop may be curtailed by a CYCLE statement. If multiple loops are associated to the loop construct, there must be no branches to any of the loop termination statements except for the innermost associated loop.

Example-9

```
#pragma omp parallel shared(a,b,c)
{
    for (i=0; i<M; i++)
        for(j=0; j<N; j++)
            for (k=0; k<P; k++)
                c[i][k] += a[i][j] * b[j][k];
}
```

Only the outer loop (i) is parallel.
Loops on j,k are serial.
Possible solutions:
 (i) change it to 1D array
 (ii) nested parallelism

Example-9

```
#pragma omp parallel shared(a,b,c)
// i,j,k are automatically private
{
    #pragma omp for collapse(3)
    for (i=0; i<M; i++)
        for(j=0; j<N; j++)
            for (k=0; k<P; k++)
                c[i][k] += a[i][j] * b[j][k];
}
```

Now everything is parallel.
c is shared. Possibly will be modified by other threads!

Example-9

```
#pragma omp parallel shared(a,b,c)
// i,j,k are automatically private
{
    #pragma omp for collapse(2)
    for (i=0; i<M; i++)
        for (k=0; k<P; k++)
            for(j=0; j<N; j++)
                c[i][k] += a[i][j] * b[j][k];
}
```

Parallelize as much as possible. Guarantee correctness and without using CRITICAL !!

Work Out - 3

- Can you rewrite the matrix multiplication using collapse?

Example 10: Random Number

```
# include <stdlib.h>
# include <stdio.h>
# include <math.h>
# include <omp.h>
# include <time.h>

void monte_carlo ( int n, int *seed );
double random_value ( int *seed );
void timestamp ( void );

void main ( ) {
    int n;
    int seed;
    timestamp ( );
    printf ( " Number of processors available = %d\n", omp_get_num_procs ( ) );
    printf ( " Number of threads =          %d\n", omp_get_max_threads ( ) );

    n = 100;
    seed = 123456789;
    monte_carlo ( n, &seed );
    printf ( "RANDOM_OPENMP\n" );
    printf ( " Normal end of execution.\n" );
    timestamp ( );
}
```

Example 10: Random Number

```
void monte_carlo ( int n, int *seed )
{
    int i, my_id, my_seed;
    double *x;
    x = ( double * ) malloc ( n * sizeof ( double ) );
    # pragma omp master
    {
        printf ( " Thread Seed I X(I)\n" );
    }
    # pragma omp parallel private ( i, my_id, my_seed ) shared ( n, x )
    {
        my_id = omp_get_thread_num ( );
        my_seed = *seed + my_id;
        printf ( " %6d %12d\n", my_id, my_seed );

        # pragma omp for
        for ( i = 0; i < n; i++ ) {
            x[i] = random_value ( &my_seed );
            printf ( " %6d %12d %6d %14.6g\n", my_id, my_seed, i, x[i] );
        }
    }
    free ( x );
    return;
}
```

Example 10: Random Number

```
double random_value ( int *seed )
{
    double r;

    *seed = ( *seed % 65536 );
    *seed = ( ( 3125 * *seed ) % 65536 );
    r = ( double ) ( *seed ) / 65536.0;

    return r;
}
```

Example 10: Random Number

```
void timestamp ( void )
{
    # define TIME_SIZE 40

    static char time_buffer[TIME_SIZE];
    const struct tm *tm;
    size_t len;
    time_t now;

    now = time ( NULL );
    tm = localtime ( &now );

    len = strftime ( time_buffer, TIME_SIZE, "%d %B %Y %l:%M:%S %p", tm );

    printf ( "%s\n", time_buffer );

    return;
    # undef TIME_SIZE
}
```

Lock Functions

- `omp_init_lock`
 - `omp_destroy_lock`
 - `omp_set_lock`
 - `omp_unset_lock`
 - `omp_test_lock`
-
- `omp_init_nest_lock`
 - `omp_destroy_nest_lock`
 - `omp_set_nest_lock`
 - `omp_unset_nest_lock`
 - `omp_test_nest_lock`